

**1(a)**

When  $0.5 - 0.3x_{i1} - 2.7x_{i2} + 3.9x_{i3} \geq \log \frac{0.7}{1-0.7}$ , that is to say when  $0.5 - 0.3x_{i1} - 2.7x_{i2} + 3.9x_{i3} \geq 0.8472979$ , the model will predict  $Y_i = 1$ .

If  $X_i = (3,1,1)$ ,  $0.5 - 0.3x_{i1} - 2.7x_{i2} + 3.9x_{i3} = 0.5 - 0.3 * 3 - 2.7 * 1 + 3.9 * 1 = 0.8$ , which is smaller than 0.8472979, so the model will predict  $Y_i = 0$ .

If  $t$  is smaller, false-positive rate will increase and false-negative rate will decrease.

**1(b)**

For  $Prob[Y = k | \hat{Y} = k] = p_k$ , the 0-1 loss for  $\hat{Y} = k$  ( $l_k$ ), is  $1 - p_k$ .

$k$	0	1	2	3	4
$l_k$	0.8	0.65	0.85	0.8	0.9

➤ #2(a)

```
> library(readr)
> wine <- read_csv("wine.csv")
Parsed with column specification:
cols(
  fixed.acidity = col_double(),
  volatile.acidity = col_double(),
  citric.acid = col_double(),
  residual.sugar = col_double(),
  chlorides = col_double(),
  free.sulfur.dioxide = col_double(),
  total.sulfur.dioxide = col_double(),
  density = col_double(),
  pH = col_double(),
  sulphates = col_double(),
  alcohol = col_double(),
  quality = col_double()
)
> view(wine)

> # Data pre-processing

> #set a high_quality column as Y
> wine$high_quality = 0
> wine$high_quality[wine$quality > 5] = 1
> wine_high_qulaity = subset(wine, select =
c("fixed.acidity", "volatile.acidity", "citric.acid", "residu
al.sugar",
+
"chlorides", "free.sulfur.dioxide", "total.sulfur.dioxide",
+
"density", "pH", "sulphates", "alcohol", "high_quality"))

> # Random sampling into training, validation, and testing
> set.seed(15)
> training.rows <- sample(1:nrow(wine_high_qulaity),
nrow(wine_high_qulaity)*0.6)
> wine_quality_training =
wine_high_qulaity[training.rows,]
> wine_quality_vt = wine_high_qulaity[-training.rows,]

> validation.rows <-
sample(1:nrow(wine_quality_vt), nrow(wine_quality_vt)*0.5)
> wine_quality_validation =
wine_quality_vt[validation.rows,]
```

```
> wine_quality_testing = wine_quality_vt[-  
validation.rows,]
```

```
> # Check whether the rows of training+validation+testing  
equals to the total number  
> nrow(wine_high_quality) ==  
nrow(wine_quality_training)+nrow(wine_quality_validation)+  
+   nrow(wine_quality_testing)  
[1] TRUE
```

Fist, split the training set(60%) from the whole data frame, and then  
equally split the rest of the data into validation(20%) and testing  
set(20%).

➤ (i) **Logistic Regression Model**

```
> # Basic logistic regression model trained on training  
set  
> # glm:generalize linear model; family=binomial:use  
logistic regression  
> model_lrg = glm(high_quality ~  
fixed.acidity+volatile.acidity+citric.acid+residual.sugar+  
+  
chlorides+free.sulfur.dioxide+total.sulfur.dioxide+  
+           density+pH+sulphates+alcohol,  
+           data=wine_quality_training,family=binomial)
```

```

> summary(model_lrg)
Call:
glm(formula = high_quality ~ fixed.acidity + volatile.acidity +
    citric.acid + residual.sugar + chlorides + free.sulfur.dioxide +
    total.sulfur.dioxide + density + pH + sulphates + alcohol,
    family = binomial, data = wine_quality_training)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.6061  -0.8978   0.4431   0.8053   2.5536

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  3.809e+02  9.481e+01   4.017 5.88e-05 ***
fixed.acidity  1.044e-01  9.589e-02   1.089 0.276360
volatile.acidity -6.457e+00  5.311e-01 -12.157 < 2e-16 ***
citric.acid    2.417e-01  3.937e-01   0.614 0.539257
residual.sugar  2.187e-01  3.579e-02   6.112 9.83e-10 ***
chlorides      2.110e+00  2.089e+00   1.010 0.312403
free.sulfur.dioxide 6.839e-03  3.538e-03   1.933 0.053222 .
total.sulfur.dioxide 3.348e-05  1.574e-03   0.021 0.983030
density       -3.948e+02  9.613e+01  -4.107 4.01e-05 ***
pH            1.471e+00  4.769e-01   3.084 0.002040 **
sulphates     1.518e+00  4.512e-01   3.365 0.000767 ***
alcohol       5.973e-01  1.245e-01   4.799 1.60e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 3738.3  on 2937  degrees of freedom
Residual deviance: 2961.2  on 2926  degrees of freedom
AIC: 2985.2

Number of Fisher Scoring iterations: 5

> # Make predictions on the validation set
> # The type="response" option tells R to output
probabilities of the form  $P(Y = 1|X)$ , as opposed to other
information such as the logit.
> predictVal_lrg = predict(model_lrg, type="response",
newdata = wine_quality_validation)

> # Analyze predictions
> summary(predictVal_lrg)
    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.004204 0.477606 0.707180 0.662999 0.869147 0.991805

> # Confusion matrix for threshold of 0.5
> table(wine_quality_validation$high_quality,
predictVal_lrg > 0.5)

```

```

      FALSE TRUE
0    159  176
1     64  581
> # Calculate the overall accuracy
> accuracy_lrg = (159+581)/(159+176+64+581)
> accuracy_lrg
[1] 0.755102

```

The overall accuracy of logistic regression model is 75.51%.

➤ (ii) k-NN model

```

> # Create the training data set
> # Scaling: (x-mean(x))/sd(x)
> f_a = (wine_quality_training$fixed.acidity-
mean(wine_quality_training$fixed.acidity))/sd(wine_quality_
_training$fixed.acidity)
> v_a = (wine_quality_training$volatile.acidity-
mean(wine_quality_training$volatile.acidity))/sd(wine_qual
ity_training$volatile.acidity)
> c_a = (wine_quality_training$citric.acid-
mean(wine_quality_training$citric.acid))/sd(wine_quality_t
raining$citric.acid)
> r_s = (wine_quality_training$residual.sugar-
mean(wine_quality_training$residual.sugar))/sd(wine_qualit
y_training$residual.sugar)
> chl = (wine_quality_training$chlorides-
mean(wine_quality_training$chlorides))/sd(wine_quality_tra
ining$chlorides)
> f_s_d = (wine_quality_training$free.sulfur.dioxide-
mean(wine_quality_training$free.sulfur.dioxide))/sd(wine_q
uality_training$free.sulfur.dioxide)
> t_s_d = (wine_quality_training$total.sulfur.dioxide-
mean(wine_quality_training$total.sulfur.dioxide))/sd(wine_
quality_training$total.sulfur.dioxide)
> den = (wine_quality_training$density-
mean(wine_quality_training$density))/sd(wine_quality_train
ing$density)
> ph = (wine_quality_training$pH-
mean(wine_quality_training$pH))/sd(wine_quality_training$p
H)
> sul = (wine_quality_training$sulphates-
mean(wine_quality_training$sulphates))/sd(wine_quality_tra
ining$sulphates)

```

```
> alc = (wine_quality_training$alcohol -  
mean(wine_quality_training$alcohol))/sd(wine_quality_train  
ing$alcohol)  
> h_q = as.factor(wine_quality_training$high_quality)
```

```
knn_train=data.frame(f_a,v_a,c_a,r_s,chl,f_s_d,t_s_d,den,p  
h,sul,alc)  
knn_training=data.frame(f_a,v_a,c_a,r_s,chl,f_s_d,t_s_d,de  
n,ph,sul,alc,h_q)
```

```
> # Create the validation data set  
> # Scaling:(x-mean(x))/sd(x)  
> f_a = (wine_quality_validation$fixed.acidity -  
mean(wine_quality_validation$fixed.acidity))/sd(wine_quali  
ty_validation$fixed.acidity)  
> v_a = (wine_quality_validation$volatile.acidity -  
mean(wine_quality_validation$volatile.acidity))/sd(wine_qu  
ality_validation$volatile.acidity)  
> c_a = (wine_quality_validation$citric.acid -  
mean(wine_quality_validation$citric.acid))/sd(wine_quality  
_validation$citric.acid)  
> r_s = (wine_quality_validation$residual.sugar -  
mean(wine_quality_validation$residual.sugar))/sd(wine_qual  
ity_validation$residual.sugar)  
> chl = (wine_quality_validation$chlorides -  
mean(wine_quality_validation$chlorides))/sd(wine_quality_v  
alidation$chlorides)  
> f_s_d = (wine_quality_validation$free.sulfur.dioxide -  
mean(wine_quality_validation$free.sulfur.dioxide))/sd(wine  
_quality_validation$free.sulfur.dioxide)  
> t_s_d = (wine_quality_validation$total.sulfur.dioxide -  
mean(wine_quality_validation$total.sulfur.dioxide))/sd(win  
e_quality_validation$total.sulfur.dioxide)  
> den = (wine_quality_validation$density -  
mean(wine_quality_validation$density))/sd(wine_quality_val  
idation$density)  
> ph = (wine_quality_validation$pH -  
mean(wine_quality_validation$pH))/sd(wine_quality_validati  
on$pH)  
> sul = (wine_quality_validation$sulphates -  
mean(wine_quality_validation$sulphates))/sd(wine_quality_v  
alidation$sulphates)  
> alc = (wine_quality_validation$alcohol -  
mean(wine_quality_validation$alcohol))/sd(wine_quality_val  
idation$alcohol)  
> h_q = as.factor(wine_quality_validation$high_quality)
```

```

knn_valid=data.frame(f_a,v_a,c_a,r_s,ch1,f_s_d,t_s_d,den,p
h,sul,alc)
knn_validation=data.frame(f_a,v_a,c_a,r_s,ch1,f_s_d,t_s_d,
den,ph,sul,alc,h_q)
>
> # Test the overall accuracy on the testing set.
> accuracy_knn = function(actual, predicted) {
+   return(mean(actual == predicted))
+ }

> # load the package 'class'
> library(class)

> #Make predictions on validation set using knn
> predicted_validation = knn(train = knn_train, test =
knn_valid, cl = knn_training$h_q, k = 5)
> high_quality_validation = knn_validation$h_q
>
> #Confusion Matrix
> table(high_quality_validation, predicted_validation)
               predicted_validation
high_quality_validation  0    1
                        0 205 130
                        1   99 546

> #Calculate the accuracy of knn
> accuracy_knn(actual = high_quality_validation,predicted
= predicted_validation)
[1] 0.7663265

```

The overall accuracy of logistic regression model is 76.63%. The performance of Logistic Regression Model is better than that of KNN Model. So, I choose *Logistic Regression Model*.

```

> # Choose KNN model
> # Create the testing data set
> # Scaling:(x-mean(x))/sd(x)
> f_a = (wine_quality_testing$fixed.acidity-
mean(wine_quality_testing$fixed.acidity))/sd(wine_quality_
testing$fixed.acidity)
> v_a = (wine_quality_testing$volatile.acidity-
mean(wine_quality_testing$volatile.acidity))/sd(wine_quali
ty_testing$volatile.acidity)
> c_a = (wine_quality_testing$citric.acid-
mean(wine_quality_testing$citric.acid))/sd(wine_quality_te
sting$citric.acid)

```

```

> r_s = (wine_quality_testing$residual.sugar-
mean(wine_quality_testing$residual.sugar))/sd(wine_quality_
_testing$residual.sugar)
> chl = (wine_quality_testing$chlorides-
mean(wine_quality_testing$chlorides))/sd(wine_quality_test
ing$chlorides)
> f_s_d = (wine_quality_testing$free.sulfur.dioxide-
mean(wine_quality_testing$free.sulfur.dioxide))/sd(wine_qu
ality_testing$free.sulfur.dioxide)
> t_s_d = (wine_quality_testing$total.sulfur.dioxide-
mean(wine_quality_testing$total.sulfur.dioxide))/sd(wine_q
uality_testing$total.sulfur.dioxide)
> den = (wine_quality_testing$density-
mean(wine_quality_testing$density))/sd(wine_quality_testin
g$density)
> ph = (wine_quality_testing$pH-
mean(wine_quality_testing$pH))/sd(wine_quality_testing$pH)
> sul = (wine_quality_testing$sulphates-
mean(wine_quality_testing$sulphates))/sd(wine_quality_test
ing$sulphates)
> alc = (wine_quality_testing$alcohol-
mean(wine_quality_testing$alcohol))/sd(wine_quality_testin
g$alcohol)
> h_q = as.factor(wine_quality_testing$high_quality)
>
knn_test=data.frame(f_a,v_a,c_a,r_s,chl,f_s_d,t_s_d,den,ph
,sul,alc)
>
knn_testing=data.frame(f_a,v_a,c_a,r_s,chl,f_s_d,t_s_d,den
,ph,sul,alc,h_q)
> #Make predictions on testing set using knn
> predicted_testing = knn(train = knn_valid, test =
knn_test, cl = knn_validation$h_q, k = 5)
> high_quality_testing = knn_testing$h_q
> #Calculate the accuracy of knn
> accuracy_knn(actual = high_quality_testing,predicted =
predicted_testing)
[1] 0.7408163

```

So, the testing error of the KNN Model is 0. 2591837 (1-0.7408163).

## ➤ 2(b)

```

> # Re-split the testing set
> set.seed(8)
> training.rows <- sample(1:nrow(wine_high_qulaity),
nrow(wine_high_qulaity)*0.7)

```



```
> wine_quality_training =
wine_high_quality[training.rows,]
> wine_quality_testing = wine_high_quality[-
training.rows,]
```

```
> #Cross Validation for K-nn Model
```

```
> #Load necessary packages for Cross Validation
```

```
> library(caret)
```

```
> library(e1071)
```

```
> # Create the training data set
```

```
> # Scaling:  $(x - \text{mean}(x)) / \text{sd}(x)$ 
```

```
> f_a = (wine_quality_training$fixed.acidity -
mean(wine_quality_training$fixed.acidity)) / sd(wine_quality_
_training$fixed.acidity)
```

```
> v_a = (wine_quality_training$volatile.acidity -
mean(wine_quality_training$volatile.acidity)) / sd(wine_qual
ity_training$volatile.acidity)
```

```
> c_a = (wine_quality_training$citric.acid -
mean(wine_quality_training$citric.acid)) / sd(wine_quality_t
raining$citric.acid)
```

```
> r_s = (wine_quality_training$residual.sugar -
mean(wine_quality_training$residual.sugar)) / sd(wine_qualit
y_training$residual.sugar)
```

```
> chl = (wine_quality_training$chlorides -
mean(wine_quality_training$chlorides)) / sd(wine_quality_tra
ining$chlorides)
```

```
> f_s_d = (wine_quality_training$free.sulfur.dioxide -
mean(wine_quality_training$free.sulfur.dioxide)) / sd(wine_q
uality_training$free.sulfur.dioxide)
```

```
> t_s_d = (wine_quality_training$total.sulfur.dioxide -
mean(wine_quality_training$total.sulfur.dioxide)) / sd(wine_
quality_training$total.sulfur.dioxide)
```

```
> den = (wine_quality_training$density -
mean(wine_quality_training$density)) / sd(wine_quality_train
ing$density)
```

```
> ph = (wine_quality_training$pH -
mean(wine_quality_training$pH)) / sd(wine_quality_training$p
H)
```

```
> sul = (wine_quality_training$sulphates -
mean(wine_quality_training$sulphates)) / sd(wine_quality_tra
ining$sulphates)
```

```
> alc = (wine_quality_training$alcohol -
mean(wine_quality_training$alcohol)) / sd(wine_quality_train
ing$alcohol)
```

```

> h_q = as.factor(wine_quality_training$high_quality)

knn_train=data.frame(f_a,v_a,c_a,r_s,chl,f_s_d,t_s_d,den,p
h,sul,alc)
knn_training=data.frame(f_a,v_a,c_a,r_s,chl,f_s_d,t_s_d,de
n,ph,sul,alc,h_q)

> # Create the testing data set
> # Scaling:(x-mean(x))/sd(x)
> f_a = (wine_quality_testing$fixed.acidity-
mean(wine_quality_testing$fixed.acidity))/sd(wine_quality_
testing$fixed.acidity)
> v_a = (wine_quality_testing$volatile.acidity-
mean(wine_quality_testing$volatile.acidity))/sd(wine_quali
ty_testing$volatile.acidity)
> c_a = (wine_quality_testing$citric.acid-
mean(wine_quality_testing$citric.acid))/sd(wine_quality_te
sting$citric.acid)
> r_s = (wine_quality_testing$residual.sugar-
mean(wine_quality_testing$residual.sugar))/sd(wine_quality
_testing$residual.sugar)
> chl = (wine_quality_testing$chlorides-
mean(wine_quality_testing$chlorides))/sd(wine_quality_test
ing$chlorides)
> f_s_d = (wine_quality_testing$free.sulfur.dioxide-
mean(wine_quality_testing$free.sulfur.dioxide))/sd(wine_qu
ality_testing$free.sulfur.dioxide)
> t_s_d = (wine_quality_testing$total.sulfur.dioxide-
mean(wine_quality_testing$total.sulfur.dioxide))/sd(wine_q
uality_testing$total.sulfur.dioxide)
> den = (wine_quality_testing$density-
mean(wine_quality_testing$density))/sd(wine_quality_testin
g$density)
> ph = (wine_quality_testing$pH-
mean(wine_quality_testing$pH))/sd(wine_quality_testing$pH)
> sul = (wine_quality_testing$sulphates-
mean(wine_quality_testing$sulphates))/sd(wine_quality_test
ing$sulphates)
> alc = (wine_quality_testing$alcohol-
mean(wine_quality_testing$alcohol))/sd(wine_quality_testin
g$alcohol)
> h_q = as.factor(wine_quality_testing$high_quality)

knn_test=data.frame(f_a,v_a,c_a,r_s,chl,f_s_d,t_s_d,den,ph
,sul,alc)
knn_testing=data.frame(f_a,v_a,c_a,r_s,chl,f_s_d,t_s_d,den
,ph,sul,alc,h_q)

```

```
> # 6-fold cross validation
> trControl <- trainControl(method = "cv", number = 6)

> fit_cv <- train(h_q ~ ., method = "knn", tuneGrid =
expand.grid(k = 1:10), metric="Accuracy",
+           trControl = trControl, data=knn_training)

> fit_cv
```

k-Nearest Neighbors

```
3428 samples
 11 predictor
  2 classes: '0', '1'
```

No pre-processing

Resampling: Cross-Validated (6 fold)

Summary of sample sizes: 2857, 2858, 2856, 2856, 2857, 2856, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
1	0.7783040	0.5023447
2	0.7354177	0.4042998
3	0.7514684	0.4318507
4	0.7415386	0.4068175
5	0.7494145	0.4220047
6	0.7494145	0.4206332
7	0.7555517	0.4302209
8	0.7520409	0.4228866
9	0.7605051	0.4403283
10	0.7575903	0.4339997

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was k = 1.

```
> # The best model is k=1.
> #Make predictions on testing set using kNN
> predicted_testing = knn(train = knn_train, test =
knn_test, cl = knn_training$h_q, k = 1)
> high_quality_testing = knn_testing$h_q
>
> #Calculate the accuracy of kNN
> accuracy_knn(actual = high_quality_testing, predicted =
predicted_testing)
[1] 0.7959184
```

Due to its highest overall accuracy, **0.7959184**, I choose k=1 as the best model, and its testing error is **0.2040816** (1-0.7959184).